

Big_Data_Analytics_Project_Data_Transformation_&_Modeling_Mar16

March 16, 2025

```
[74]: import pandas as pd
```

```
url = "https://archive.ics.uci.edu/static/public/350/data.csv"
data = pd.read_csv(url, sep= ',')
print(data.head())
```

	ID	X1	X2	X3	X4	X5	X6	X7	X8	X9	...	X15	X16	X17	X18	\
0	1	20000	2	2	1	24	2	2	-1	-1	...	0	0	0	0	
1	2	120000	2	2	2	26	-1	2	0	0	...	3272	3455	3261	0	
2	3	90000	2	2	2	34	0	0	0	0	...	14331	14948	15549	1518	
3	4	50000	2	2	1	37	0	0	0	0	...	28314	28959	29547	2000	
4	5	50000	1	2	1	57	-1	0	-1	0	...	20940	19146	19131	2000	

	X19	X20	X21	X22	X23	Y
0	689	0	0	0	0	1
1	1000	1000	1000	0	2000	1
2	1500	1000	1000	1000	5000	0
3	2019	1200	1100	1069	1000	0
4	36681	10000	9000	689	679	0

[5 rows x 25 columns]

```
[75]: #Renaming columns
```

```
data.rename(columns={'X1': 'LIMIT_BAL', 'X2': 'SEX', 'X3': 'EDUCATION', 'X4': 'MARRIAGE', 'X5': 'AGE', 'X6': 'PAY_0', 'X7': 'PAY_2', 'X8': 'PAY_3', 'X9': 'PAY_4', 'X10': 'PAY_5', 'X11': 'PAY_6', 'X12': 'BILL_AMT1', 'X13': 'BILL_AMT2', 'X14': 'BILL_AMT3', 'X15': 'BILL_AMT4', 'X16': 'BILL_AMT5', 'X17': 'BILL_AMT6', 'X18': 'PAY_AMT1', 'X19': 'PAY_AMT2', 'X20': 'PAY_AMT3', 'X21': 'PAY_AMT4', 'X22': 'PAY_AMT5', 'X23': 'PAY_AMT6'}, inplace=True)
data.head()
```

```
[75]:
```

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	\
0	1	20000	2			24	2	2	-1	-1	
1	2	120000	2			26	-1	2	0	0	
2	3	90000	2			34	0	0	0	0	
3	4	50000	2		1	37	0	0	0	0	

	4	5	50000	1	2	1	57	-1	0	-1	0
...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	\				
0	...	0	0	0	0	689	0				
1	...	3272	3455	3261	0	1000	1000				
2	...	14331	14948	15549	1518	1500	1000				
3	...	28314	28959	29547	2000	2019	1200				
4	...	20940	19146	19131	2000	36681	10000				

	PAY_AMT4	PAY_AMT5	PAY_AMT6	Y
0	0	0	0	1
1	1000	0	2000	1
2	1000	1000	5000	0
3	1100	1069	1000	0
4	9000	689	679	0

[5 rows x 25 columns]

```
[76]: # Replacing education values = 0, 5 and 6 with 4, since 0, 5 and 6 are not
      ↪ defined

fill = (data.EDUCATION == 0) | (data.EDUCATION == 5) | (data.EDUCATION == 6)
data.loc[fill, 'EDUCATION'] = 4

print('EDUCATION ' + str(sorted(data['EDUCATION'].unique())))
```

EDUCATION [1, 2, 3, 4]

```
[77]: # Replacing marital status value = 0 to 3, since 0 is not defined

fill = (data.MARRIAGE == 0)
data.loc[fill, 'MARRIAGE'] = 3

print('MARRIAGE ' + str(sorted(data['MARRIAGE'].unique())))
```

MARRIAGE [1, 2, 3]

```
[78]: # Applying One-Hot Encoding technique to categorical variables

from sklearn.preprocessing import OneHotEncoder

categorical_variables = ['SEX', 'EDUCATION', 'MARRIAGE', 'PAY_0', 'PAY_2',
      ↪ 'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6']

encoder = OneHotEncoder(sparse_output=False)
one_hot_encoded = encoder.fit_transform(data[categorical_variables])
```

```

one_hot_df = pd.DataFrame(one_hot_encoded, columns=encoder.
    ↪get_feature_names_out(categorical_variables))

df_encoded = pd.concat([data, one_hot_df], axis=1)

df_encoded = df_encoded.drop(categorical_variables, axis=1)
print(f"Encoded Employee data : \n{df_encoded}")

```

Encoded Employee data :

	ID	LIMIT_BAL	AGE	BILL_AMT1	BILL_AMT2	BILL_AMT3	BILL_AMT4	\
0	1	20000	24	3913	3102	689	0	
1	2	120000	26	2682	1725	2682	3272	
2	3	90000	34	29239	14027	13559	14331	
3	4	50000	37	46990	48233	49291	28314	
4	5	50000	57	8617	5670	35835	20940	
...	
29995	29996	220000	39	188948	192815	208365	88004	
29996	29997	150000	43	1683	1828	3502	8979	
29997	29998	30000	37	3565	3356	2758	20878	
29998	29999	80000	41	-1645	78379	76304	52774	
29999	30000	50000	46	47929	48905	49764	36535	

	BILL_AMT5	BILL_AMT6	PAY_AMT1	...	PAY_6_-2	PAY_6_-1	PAY_6_0	\
0	0	0	0	...	1.0	0.0	0.0	
1	3455	3261	0	...	0.0	0.0	0.0	
2	14948	15549	1518	...	0.0	0.0	1.0	
3	28959	29547	2000	...	0.0	0.0	1.0	
4	19146	19131	2000	...	0.0	0.0	1.0	
...	
29995	31237	15980	8500	...	0.0	0.0	1.0	
29996	5190	0	1837	...	0.0	0.0	1.0	
29997	20582	19357	0	...	0.0	0.0	1.0	
29998	11855	48944	85900	...	0.0	1.0	0.0	
29999	32428	15313	2078	...	0.0	0.0	1.0	

	PAY_6_2	PAY_6_3	PAY_6_4	PAY_6_5	PAY_6_6	PAY_6_7	PAY_6_8
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	1.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0
...
29995	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29996	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29997	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29998	0.0	0.0	0.0	0.0	0.0	0.0	0.0
29999	0.0	0.0	0.0	0.0	0.0	0.0	0.0

[30000 rows x 89 columns]

```
[79]: numeric_variables = ['LIMIT_BAL', 'AGE',  
    ↪ 'BILL_AMT1', 'BILL_AMT2', 'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6',  
    ↪ 'PAY_AMT1', 'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6']  
  
#Calculating and printing the Pearson correlation matrix  
print("Pearson Correlation Matrix of numeric variables:")  
pearson_correlation_matrix = data[numeric_variables].corr().round(2)  
pearson_correlation_matrix
```

Pearson Correlation Matrix of numeric variables:

```
[79]:
```

	LIMIT_BAL	AGE	BILL_AMT1	BILL_AMT2	BILL_AMT3	BILL_AMT4	\
LIMIT_BAL	1.00	0.14	0.29	0.28	0.28	0.29	
AGE	0.14	1.00	0.06	0.05	0.05	0.05	
BILL_AMT1	0.29	0.06	1.00	0.95	0.89	0.86	
BILL_AMT2	0.28	0.05	0.95	1.00	0.93	0.89	
BILL_AMT3	0.28	0.05	0.89	0.93	1.00	0.92	
BILL_AMT4	0.29	0.05	0.86	0.89	0.92	1.00	
BILL_AMT5	0.30	0.05	0.83	0.86	0.88	0.94	
BILL_AMT6	0.29	0.05	0.80	0.83	0.85	0.90	
PAY_AMT1	0.20	0.03	0.14	0.28	0.24	0.23	
PAY_AMT2	0.18	0.02	0.10	0.10	0.32	0.21	
PAY_AMT3	0.21	0.03	0.16	0.15	0.13	0.30	
PAY_AMT4	0.20	0.02	0.16	0.15	0.14	0.13	
PAY_AMT5	0.22	0.02	0.17	0.16	0.18	0.16	
PAY_AMT6	0.22	0.02	0.18	0.17	0.18	0.18	

	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	PAY_AMT4	\
LIMIT_BAL	0.30	0.29	0.20	0.18	0.21	0.20	
AGE	0.05	0.05	0.03	0.02	0.03	0.02	
BILL_AMT1	0.83	0.80	0.14	0.10	0.16	0.16	
BILL_AMT2	0.86	0.83	0.28	0.10	0.15	0.15	
BILL_AMT3	0.88	0.85	0.24	0.32	0.13	0.14	
BILL_AMT4	0.94	0.90	0.23	0.21	0.30	0.13	
BILL_AMT5	1.00	0.95	0.22	0.18	0.25	0.29	
BILL_AMT6	0.95	1.00	0.20	0.17	0.23	0.25	
PAY_AMT1	0.22	0.20	1.00	0.29	0.25	0.20	
PAY_AMT2	0.18	0.17	0.29	1.00	0.24	0.18	
PAY_AMT3	0.25	0.23	0.25	0.24	1.00	0.22	
PAY_AMT4	0.29	0.25	0.20	0.18	0.22	1.00	
PAY_AMT5	0.14	0.31	0.15	0.18	0.16	0.15	
PAY_AMT6	0.16	0.12	0.19	0.16	0.16	0.16	

PAY_AMT5 PAY_AMT6

LIMIT_BAL	0.22	0.22
AGE	0.02	0.02
BILL_AMT1	0.17	0.18
BILL_AMT2	0.16	0.17
BILL_AMT3	0.18	0.18
BILL_AMT4	0.16	0.18
BILL_AMT5	0.14	0.16
BILL_AMT6	0.31	0.12
PAY_AMT1	0.15	0.19
PAY_AMT2	0.18	0.16
PAY_AMT3	0.16	0.16
PAY_AMT4	0.15	0.16
PAY_AMT5	1.00	0.15
PAY_AMT6	0.15	1.00

```
[80]: # Printing summary statistics of numeric variables before replacing outliers
      ↪with median:
print("\nSummary Statistics of numeric variables:")
data[numeric_variables].describe().transpose()
```

Summary Statistics of numeric variables:

```
[80]:
```

	count	mean	std	min	25%	\
LIMIT_BAL	30000.0	167484.322667	129747.661567	10000.0	50000.00	
AGE	30000.0	35.485500	9.217904	21.0	28.00	
BILL_AMT1	30000.0	51223.330900	73635.860576	-165580.0	3558.75	
BILL_AMT2	30000.0	49179.075167	71173.768783	-69777.0	2984.75	
BILL_AMT3	30000.0	47013.154800	69349.387427	-157264.0	2666.25	
BILL_AMT4	30000.0	43262.948967	64332.856134	-170000.0	2326.75	
BILL_AMT5	30000.0	40311.400967	60797.155770	-81334.0	1763.00	
BILL_AMT6	30000.0	38871.760400	59554.107537	-339603.0	1256.00	
PAY_AMT1	30000.0	5663.580500	16563.280354	0.0	1000.00	
PAY_AMT2	30000.0	5921.163500	23040.870402	0.0	833.00	
PAY_AMT3	30000.0	5225.681500	17606.961470	0.0	390.00	
PAY_AMT4	30000.0	4826.076867	15666.159744	0.0	296.00	
PAY_AMT5	30000.0	4799.387633	15278.305679	0.0	252.50	
PAY_AMT6	30000.0	5215.502567	17777.465775	0.0	117.75	

	50%	75%	max
LIMIT_BAL	140000.0	240000.00	1000000.0
AGE	34.0	41.00	79.0
BILL_AMT1	22381.5	67091.00	964511.0
BILL_AMT2	21200.0	64006.25	983931.0
BILL_AMT3	20088.5	60164.75	1664089.0
BILL_AMT4	19052.0	54506.00	891586.0
BILL_AMT5	18104.5	50190.50	927171.0

BILL_AMT6	17071.0	49198.25	961664.0
PAY_AMT1	2100.0	5006.00	873552.0
PAY_AMT2	2009.0	5000.00	1684259.0
PAY_AMT3	1800.0	4505.00	896040.0
PAY_AMT4	1500.0	4013.25	621000.0
PAY_AMT5	1500.0	4031.50	426529.0
PAY_AMT6	1500.0	4000.00	528666.0

```
[81]: # Defining function to replace outliers with the median
def replace_outliers_with_median(data, column):
    median = data[column].median()
    q1 = data[column].quantile(0.25)
    q3 = data[column].quantile(0.75)
    iqr = q3 - q1
    lower_bound = q1 - 1.5 * iqr
    upper_bound = q3 + 1.5 * iqr

    data[column] = data[column].apply(lambda x: median if x < lower_bound or x >
    upper_bound else x)
```

```
[82]: # Applying the function to the columns with outliers
replace_outliers_with_median(data, 'LIMIT_BAL')
replace_outliers_with_median(data, 'BILL_AMT1')
replace_outliers_with_median(data, 'BILL_AMT2')
replace_outliers_with_median(data, 'BILL_AMT3')
replace_outliers_with_median(data, 'BILL_AMT4')
replace_outliers_with_median(data, 'BILL_AMT5')
replace_outliers_with_median(data, 'BILL_AMT6')
replace_outliers_with_median(data, 'PAY_AMT1')
replace_outliers_with_median(data, 'PAY_AMT2')
replace_outliers_with_median(data, 'PAY_AMT3')
replace_outliers_with_median(data, 'PAY_AMT4')
replace_outliers_with_median(data, 'PAY_AMT5')
replace_outliers_with_median(data, 'PAY_AMT6')
```

```
[83]: print("\nSummary Statistics of numeric variables after replacing outliers with
    median:")
data[numeric_variables].describe().transpose()
```

Summary Statistics of numeric variables after replacing outliers with median:

```
[83]:
```

	count	mean	std	min	25%	50%	\
LIMIT_BAL	30000.0	164824.322667	125192.989579	10000.0	50000.00	140000.0	
AGE	30000.0	35.485500	9.217904	21.0	28.00	34.0	
BILL_AMT1	30000.0	33109.792100	37794.502441	-15308.0	3563.00	22381.5	
BILL_AMT2	30000.0	31669.887567	36414.965831	-69777.0	2984.75	21198.5	

BILL_AMT3	30000.0	29736.798283	34293.746628	-61506.0	2667.75	20088.5
BILL_AMT4	30000.0	26625.608833	30764.323883	-65167.0	2329.00	19052.0
BILL_AMT5	30000.0	24247.883050	28331.916539	-61372.0	1763.75	18104.5
BILL_AMT6	30000.0	23287.670000	27946.193005	-57060.0	1259.75	17071.0
PAY_AMT1	30000.0	2681.008300	2557.378286	0.0	1000.00	2100.0
PAY_AMT2	30000.0	2586.259267	2533.473459	0.0	833.00	2009.0
PAY_AMT3	30000.0	2267.026400	2396.721279	0.0	390.00	1800.0
PAY_AMT4	30000.0	1911.001400	2056.702179	0.0	296.00	1500.0
PAY_AMT5	30000.0	1926.580500	2075.388113	0.0	252.50	1500.0
PAY_AMT6	30000.0	1893.753100	2071.970037	0.0	117.75	1500.0

	75%	max
LIMIT_BAL	240000.00	520000.0
AGE	41.00	79.0
BILL_AMT1	48707.50	162296.0
BILL_AMT2	47812.25	155508.0
BILL_AMT3	44887.75	146410.0
BILL_AMT4	37803.00	132754.0
BILL_AMT5	32030.50	122830.0
BILL_AMT6	30563.00	121062.0
PAY_AMT1	3706.00	11013.0
PAY_AMT2	3500.00	11249.0
PAY_AMT3	3005.00	10673.0
PAY_AMT4	2816.25	9584.0
PAY_AMT5	2913.50	9700.0
PAY_AMT6	2853.50	9817.0

```
[84]: #Normalizing the numeric attributes
from sklearn.preprocessing import MinMaxScaler

# Initializing the scaler
scaler = MinMaxScaler()
# Fitting and transforming the data
X = scaler.fit_transform(data[numeric_variables])

y = data['Y']
```

```
[85]: # Validating minimum and maximum values are set as 0.00 and 1.00 respectively
normalized_df_summary = pd.DataFrame(X).describe()
normalized_df_summary
```

```
[85]:
```

	0	1	2	3	4	\
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	
mean	0.303577	0.249750	0.272617	0.450305	0.438845	
std	0.245476	0.158929	0.212802	0.161640	0.164940	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.078431	0.120690	0.106253	0.322976	0.308652	

50%	0.254902	0.224138	0.212211	0.403824	0.392440
75%	0.450980	0.344828	0.360440	0.521958	0.511715
max	1.000000	1.000000	1.000000	1.000000	1.000000

	5	6	7	8	9 \
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	0.463784	0.464815	0.451082	0.243440	0.229910
std	0.155437	0.153809	0.156894	0.232214	0.225218
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.341025	0.342753	0.327415	0.090802	0.074051
50%	0.425518	0.431464	0.416181	0.190684	0.178594
75%	0.520258	0.507066	0.491927	0.336511	0.311139
max	1.000000	1.000000	1.000000	1.000000	1.000000

	10	11	12	13
count	30000.000000	30000.000000	30000.000000	30000.000000
mean	0.212408	0.199395	0.198617	0.192905
std	0.224559	0.214597	0.213958	0.211059
min	0.000000	0.000000	0.000000	0.000000
25%	0.036541	0.030885	0.026031	0.011994
50%	0.168650	0.156511	0.154639	0.152796
75%	0.281552	0.293849	0.300361	0.290669
max	1.000000	1.000000	1.000000	1.000000

```
[86]: from sklearn.model_selection import train_test_split

# Splitting data into training and testing sets: training set 70%, test set 30%
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↪random_state=42)
```

```
[87]: X_train[0:5,]
```

```
[87]: array([[0.45098039, 0.32758621, 0.08619175, 0.30972768, 0.29582139,
0.32925763, 0.33317771, 0.32034224, 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          ],
[0.07843137, 0.03448276, 0.10550438, 0.32072708, 0.30687874,
0.35350973, 0.38643446, 0.32404756, 0.23136293, 0.20632945,
0.44973297, 0.15651085, 0.06804124, 0.30355506],
[0.07843137, 0.25862069, 0.34633792, 0.51018044, 0.53209469,
0.57155633, 0.56975494, 0.51754416, 0.          , 0.41781492,
0.          , 0.2090985 , 0.36082474, 0.          ],
[0.37254902, 0.56896552, 0.70658882, 0.78763344, 0.7993228 ,
0.8473886 , 0.88366033, 0.32034224, 0.          , 0.          ,
0.          , 0.          , 0.          , 0.          ],
[0.45098039, 0.24137931, 0.09758789, 0.31863639, 0.40062333,
0.41566585, 0.40574478, 0.4475528 , 0.18314719, 0.17859365,
0.10493769, 0.15651085, 0.15463918, 0.15279617]])
```



```
[88]: from collections import Counter
      # summarizing class distribution before applying SMOTE
      print(Counter(y_train))
```

```
Counter({0: 16324, 1: 4676})
```

```
[89]: from imblearn.over_sampling import SMOTE
      # Applying SMOTE to transform the dataset
      oversample = SMOTE(sampling_strategy=0.5)

      #Fitting and applying the transform
      X_train, y_train = oversample.fit_resample(X_train, y_train)
```

```
[90]: # summarizing the new class distribution after applying SMOTE
      print(Counter(y_train))
```

```
Counter({0: 16324, 1: 8162})
```

```
[91]: from imblearn.under_sampling import RandomUnderSampler
      # Applying undersampling to transform the dataset
      undersample = RandomUnderSampler(sampling_strategy=0.5)

      #Fitting and applying the transform
      X_train, y_train = undersample.fit_resample(X_train, y_train)
```

```
[92]: # summarizing the new class distribution after applying undersampling
      print(Counter(y_train))
```

```
Counter({0: 16324, 1: 8162})
```

```
[93]: from sklearn.linear_model import LogisticRegression
      from sklearn.model_selection import cross_val_score
      # Initializing the logistic regression model as baseline model

      log_reg = LogisticRegression()
```

```
[94]: # Performing cross-validation on the training set
      cv_scores = cross_val_score(log_reg, X_train, y_train, cv=5, scoring='accuracy')
      print("Cross-Validation Accuracy Scores on Training Set:", cv_scores)
      print("Average Cross-Validation Accuracy on Training Set:", cv_scores.mean())
```

```
Cross-Validation Accuracy Scores on Training Set: [0.67109024 0.67245252
0.66836839 0.67204411 0.66714315]
Average Cross-Validation Accuracy on Training Set: 0.670219681836189
```

```
[95]: # Training the logistic regression model on the full training set
      log_reg.fit(X_train, y_train)
```

```
[95]: LogisticRegression()
```

```
[96]: # Making predictions with Logistic Regression on the test set
y_pred_test = log_reg.predict(X_test)
```

```
[97]: from sklearn.metrics import confusion_matrix, accuracy_score, recall_score, \
      precision_score, f1_score

# Calculating evaluation metrics for Logistic Regression
accuracy = accuracy_score(y_test, y_pred_test)
precision = precision_score(y_test, y_pred_test)
recall = recall_score(y_test, y_pred_test)
f1 = f1_score(y_test, y_pred_test)

conf_matrix = confusion_matrix(y_test, y_pred_test)
print("Confusion Matrix Logistic Regression:\n", conf_matrix)
print("Accuracy Logistic Regression:", accuracy)
print("Precision Logistic Regression:", precision)
print("Recall Logistic Regression:", recall)
print("F1 score Logistic Regression:", f1)
```

Confusion Matrix Logistic Regression:

```
[[6945  95]
 [1886  74]]
```

Accuracy Logistic Regression: 0.7798888888888889
Precision Logistic Regression: 0.4378698224852071
Recall Logistic Regression: 0.03775510204081633
F1 score Logistic Regression: 0.06951620479098168

```
[98]: tn, fp, fn, tp = confusion_matrix(y_test, y_pred_test).ravel()
print("True Negatives_LG (TN):", tn)
print("False Positives_LG (FP):", fp)
print("False Negatives_LG (FN):", fn)
print("True Positives_LG (TP):", tp)
```

True Negatives_LG (TN): 6945
False Positives_LG (FP): 95
False Negatives_LG (FN): 1886
True Positives_LG (TP): 74

```
[99]: from sklearn.ensemble import RandomForestClassifier
# Initializing the Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
[100]: # Training the Random Forest classifier
rf_classifier.fit(X_train, y_train)
```

```
[100]: RandomForestClassifier(random_state=42)
```

```
[101]: # Making predictions with Random Forest classifier on the test set
y_pred_test = rf_classifier.predict(X_test)
```

```
[102]: # Calculating evaluation metrics for Random Forest
accuracy = accuracy_score(y_test, y_pred_test)
precision = precision_score(y_test, y_pred_test)
recall = recall_score(y_test, y_pred_test)
f1 = f1_score(y_test, y_pred_test)

conf_matrix = confusion_matrix(y_test, y_pred_test)
print("Confusion Matrix Random Forest:\n", conf_matrix)
print("Accuracy Random Forest:", accuracy)
print("Precision Random Forest:", precision)
print("Recall Random Forest:", recall)
print("F1 score Random Forest:", f1)
```

Confusion Matrix Random Forest:

```
[[6473  567]
```

```
[1379  581]]
```

Accuracy Random Forest: 0.7837777777777778

Precision Random Forest: 0.5060975609756098

Recall Random Forest: 0.29642857142857143

F1 score Random Forest: 0.3738738738738739

```
[103]: from sklearn.ensemble import GradientBoostingClassifier
# Initializing the Gradient Boosting Classifier
gb_clf = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
    ↪max_depth=3, random_state=42)
```

```
[104]: # Training the Gradient Boosting model
gb_clf.fit(X_train, y_train)
```

```
[104]: GradientBoostingClassifier(random_state=42)
```

```
[105]: # Making predictions with Gradient Boosting classifier
y_pred_test = gb_clf.predict(X_test)
```

```
[106]: # Calculating evaluation metrics for Gradient Boosting
accuracy = accuracy_score(y_test, y_pred_test)
precision = precision_score(y_test, y_pred_test)
recall = recall_score(y_test, y_pred_test)
f1 = f1_score(y_test, y_pred_test)

conf_matrix = confusion_matrix(y_test, y_pred_test)
print("Confusion Matrix Gradient Boosting:\n", conf_matrix)
```

```
print("Accuracy Gradient Boosting:", accuracy)
print("Precision Gradient Boosting:", precision)
print("Recall Gradient Boosting:", recall)
print("F1 score Gradient Boosting:", f1)
```

Confusion Matrix Gradient Boosting:

```
[[6527  513]
```

```
[1439  521]]
```

Accuracy Gradient Boosting: 0.7831111111111111

Precision Gradient Boosting: 0.5038684719535783

Recall Gradient Boosting: 0.26581632653061227

F1 score Gradient Boosting: 0.34802939211756845

```
[ ]: !apt-get install -y texlive-xetex texlive-fonts-recommended texlive-latex-extra
!apt-get install texlive texlive-latex-extra pandoc
!jupyter nbconvert --to pdf "/content/Big_Data_Analytics_Project_Data_
Transformation & Modeling_Mar16.ipynb"
```

Reading package lists... Done

Building dependency tree... Done

Reading state information... Done

The following additional packages will be installed:

dvisvgm fonts-droid-fallback fonts-lato fonts-lmodern fonts-noto-mono fonts-texgyre

fonts-urw-base35 libapache-pom-java libcommons-logging-java libcommons-parent-java

libfontbox-java libfontenc1 libgs9 libgs9-common libidn12 libijs-0.35

libjbig2dec0 libkpathsea6

libpdfbox-java libptexenc1 libruby3.0 libsynctex2 libteckit0 libtexlua53

libtexluaajit2 libwoff1

libzip-0-13 lmodern poppler-data preview-latex-style rake ruby ruby-net-telnet ruby-rubygems

ruby-webrick ruby-xmlrpc ruby3.0 rubygems-integration t1utils teckit tex-common tex-gyre

texlive-base texlive-binaries texlive-latex-base texlive-latex-recommended texlive-pictures

texlive-plain-generic tipa xfonts-encodings xfonts-utils

Suggested packages:

fonts-noto fonts-freefont-otf | fonts-freefont-ttf libavalon-framework-java

libcommons-logging-java-doc libexcalibur-logkit-java liblog4j1.2-java poppler-utils ghostscript

fonts-japanese-mincho | fonts-ipafont-mincho fonts-japanese-gothic | fonts-ipafont-gothic

fonts-arphic-ukai fonts-arphic-uming fonts-nanum ri ruby-dev bundler debhelper gv

| postscript-viewer perl-tk xpdf | pdf-viewer xzdec texlive-fonts-recommended-doc

texlive-latex-base-doc python3-pygments icc-profiles libfile-which-perl